

Homework-03

Neural Networks - Back-propagation pass

Background

This homework teaches how to code the fundamental training algorithm for neural networks: back-prop and stochastic gradient descent. While there are a number of tools for this, it is important to understand the main operations involved in a neural network training and code your very own SGD.

Deliverables

Small write up in *Markdown* where you will also compare the Theta you hand-crafted for HW02 vs. the one you've learnt with back-propagation.

Source code in [NeuralNetwork.py](#), [logicGates.py](#) description in [README.md](#).

API ([NeuralNetwork](#))

Points: 6

```
-- create the table of matrices  $\Theta$ 
[nil] build([int] in, [int] h1, [int] h2, ..., [int] out))
-- returns  $\Theta$ (layer)
[2D FloatTensor] getLayer([int] layer)
-- feedforward pass single vector
[1D FloatTensor] forward([1D FloatTensor] input)
-- feedforward pass design matrix (should come free after transposition)
[2D FloatTensor] forward([2D FloatTensor] input)
-- back-propagation pass single target (computes  $\partial E / \partial \Theta$ )
[nil] backward([1D FloatTensor] target)
-- back-propagation pass target matrix
-- (computes the average of  $\partial E / \partial \Theta$  across seen samples)
[nil] backward([2D FloatTensor] target)
-- update parameters
[nil] updateParams([float] eta)
```

When I import your library, I should get a class with **only** five methods specified in the API. No global variables, no other helper functions (which must be local, if needed).

API ([logicGates](#))

Points: 4

```
[nil] AND.train()
[nil] OR.train()
[nil] NOT.train()
[nil] XOR.train()
[boolean] AND.forward([boolean] x, [boolean] y)
[boolean] OR.forward([boolean] x, [boolean] y)
[boolean] NOT.forward([boolean] x)
[boolean] XOR.forward([boolean] x, [boolean] y)
```

When I import your library I should see 4 classes with **only** two public methods for each logic gate class, specified in the API. No global variables, no other helper functions (which must be local, if needed).

API (Img2Num) (Preview of next homework)

```
[nil] train()
[int] forward([28x28 ByteTensor] img)
```

When I import your library I should get a class with **only** two public methods specified in the API. No global variables, no other helper functions (which must be local, if needed, like the `oneHot()` conversion function).

Instructions

1. Create a *Python* script `NeuralNetwork.py` that contains two local empty dictionaries `Theta` and `dE_dTheta` and the five functions described in the API.
2. Implement back-propagation with a *Mean Square Error* loss function.
3. Update the matrices `Theta` with `updateParams(eta)` based on the learning rate `eta` and the gradient of the error with respect of the parameters `dE_dTheta`.
4. Train the **AND**, **OR**, **NOT** and **XOR** networks using the **NeuralNetwork** API (calling `forward()`, `backward()` and `updateParams()` in a cycle) on a hand-crafted data set (Hint: use *Python's* `and`, `or`, `not` and combination of these in order to build your data on the fly). Compare the the `Thetas` with what you set manually in HW02. One may learn the new `Thetas` with the `train()` function.
5. Train the `img2num` network using the MNIST dataset using the **NeuralNetwork** API. To speed up training, you may want to use the **NeuralNetwork** API in batch mode (where you forward chunks of the design and target matrices at a time). This means that if `size(x) = n` and you have `m` examples, then `size(X) = m × n` (I know it is transpose of what you were providing in last homework) and **NeuralNetwork's** `forward()` can take `x` or `X` and **NeuralNetwork's** `backward()` can take `y` or `Y`. Output of **NeuralNetwork's** `forward()` will be `m × c` (where `m` is number of examples and `c` is number of classes/categories).

Format

Submit ZIP file containing all the deliverables with instructions to run your script using. Your folder name should be as following:

`<your_purdue_username>_HW<0X>`

For example, my folder name for this homework will look like `aabhisht_HW03`.

Bonus

Points: 2

Add one more parameter to `backward([1D DoubleTensor] target)` which becomes now `backward(target, [string] loss)`. By default `loss` is `MSE`. Nevertheless, `loss` can be `CE`, which stands for *cross-entropy*.

The cross-entropy loss function is defined as following:

$$\begin{aligned} \text{loss}(h, y) &= -\log(\exp(h[y]) / (\sum_j \exp(h[j]))) \\ &= -h[y] + \log(\sum_j \exp(h[j])) \end{aligned}$$

Suggestions/Recommendations

- You may want to create two additional attributes (dictionaries) `a` and `z` in `NeuralNetwork` class.
- Make your own tests to verify the output of the logic gate functions against the *Python's* `and`, `or`, `not` and combination of these. You should build and train the `AND`, `OR`, `NOT` and `XOR` networks using your own `NeuralNetwork` library. You will need to convert `booleans` to the integers `0` and `1` and vice versa, in order to comply with the given API.
- You will probably need to download the MNIST data set. Instruction about how to use the data set are provided here -> <https://github.com/andresy/mnist>.
- Your `img2num` network will need `view` the 2D input into a 1D Tensor so that it can be fed to the network. The labels will need a `oneHot` encoding.